

# Transparent SQL-of-Thought: Explainable Multi-Agent Text-to-SQL Generation

Advait Shinde  
Toney Zhen  
Jasper Morgal



Advait Shinde  
Professional Vibe Coder and  
Photographer



Toney Zhen  
Professional Spectator, Cheerleader,  
Benchwarmer



Jasper Paul Morgal  
Also Professional Vibe Coder

Photos courtesy of Advait Shinde  
*Taken during company bonding offsite hike*

# AGENDA

01 The Challenge

02 The Solution

03 Timeline

04 Our Implementation

05 Our Improvements

06 Results & Metrics

# The Challenge

Convert natural language queries into SQL queries

## Traditional Approaches

- ✗ Single-shot generation
- ✗ No structured reasoning
- ✗ Generic error feedback
- ✗ Blind refinement

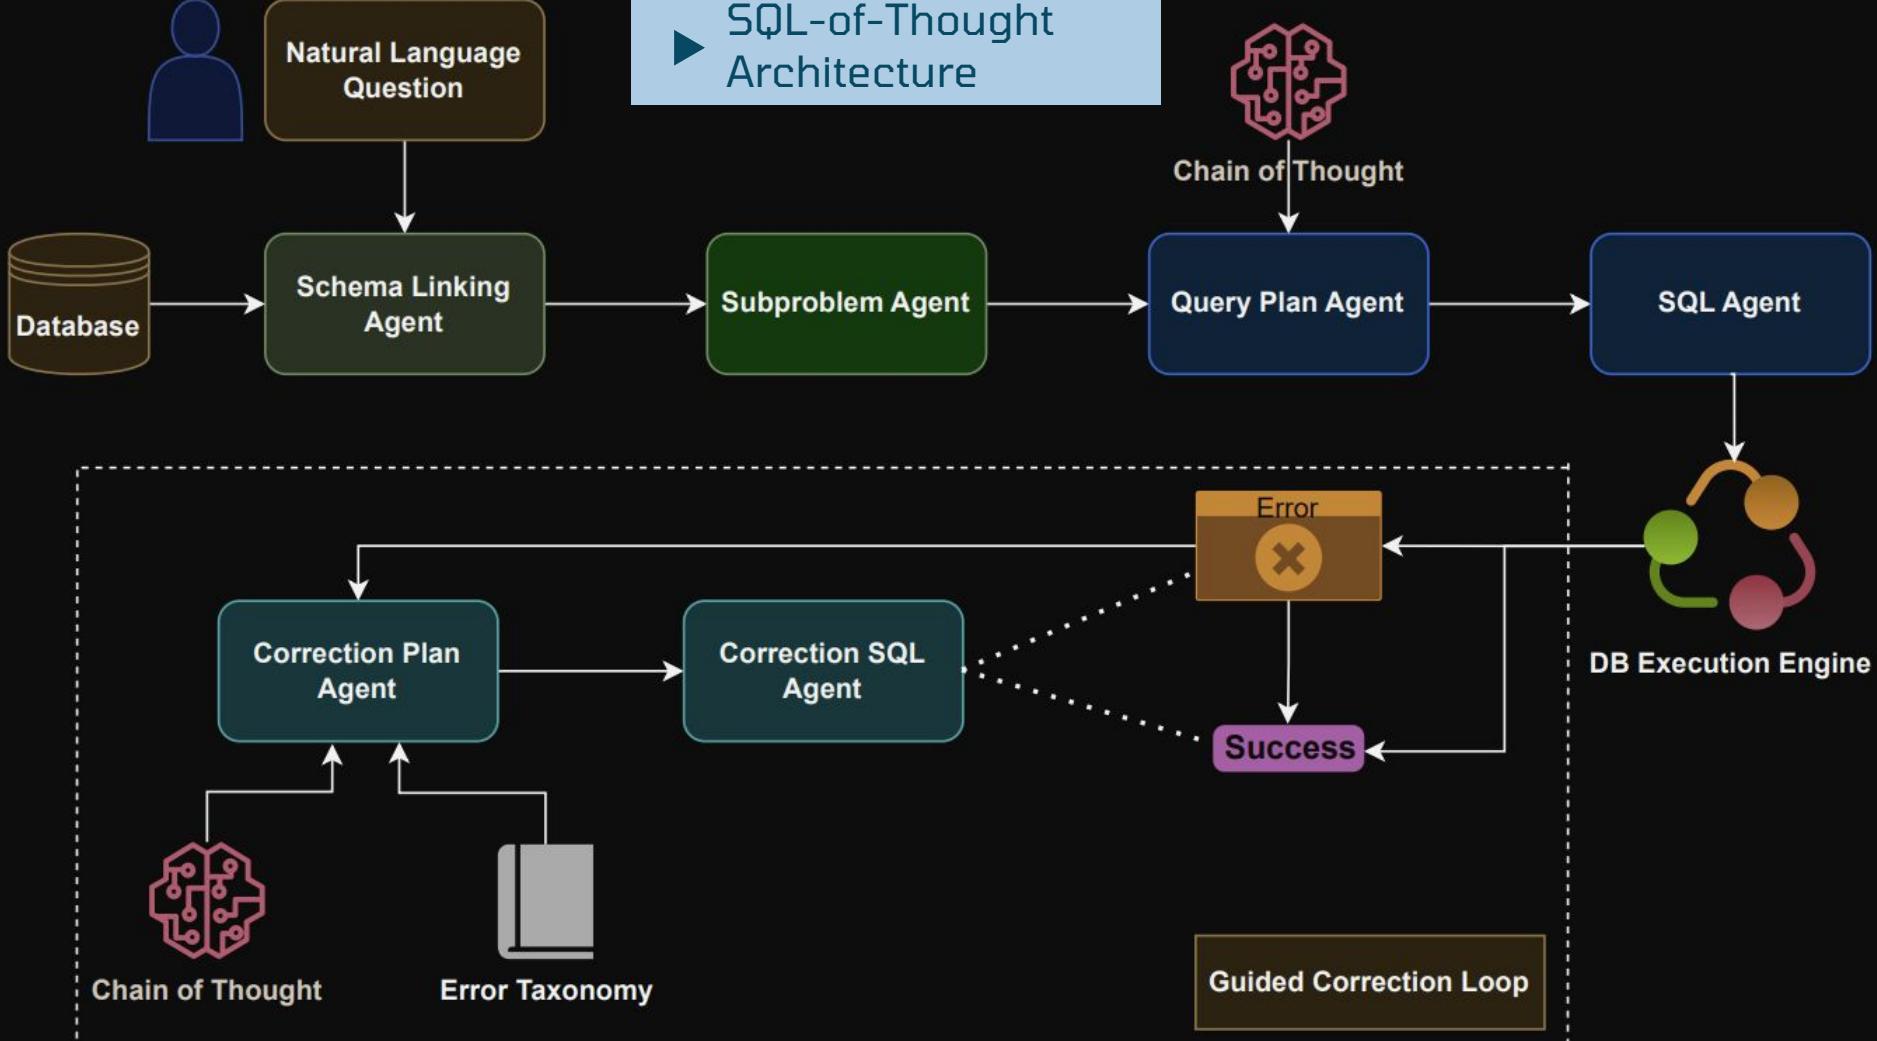
Many generated queries are syntactically valid but logically incorrect. Execution errors alone don't provide sufficient guidance for query correction.

## Real World Needs

- ✓ Complex cross-domain queries
- ✓ Multi-table joins
- ✓ Logical error detection
- ✓ Iterative refinement

Natural Language Question

## SQL-of-Thought Architecture



# GUIDED ERROR CORRECTION : 9 Categories, 31 Error Types

Syntax

sql\_syntax\_error,  
invalid\_alias

Sub  
Query

unused, missing,  
correlation\_error

Set  
Operations

union/intersect/exc  
ept missing

Value  
Errors

hardcoded\_value,  
format\_wrong

Join  
Errors

join\_missing,  
wrong\_type,  
extra\_table

Aggregation

agg\_no\_groupby,  
having\_vs\_where

Schema  
Link

table/col\_missing,  
ambiguous\_col,  
incorrect\_fk

Filter  
Errors

where\_missing,  
wrong\_col,  
type\_mismatch

Others

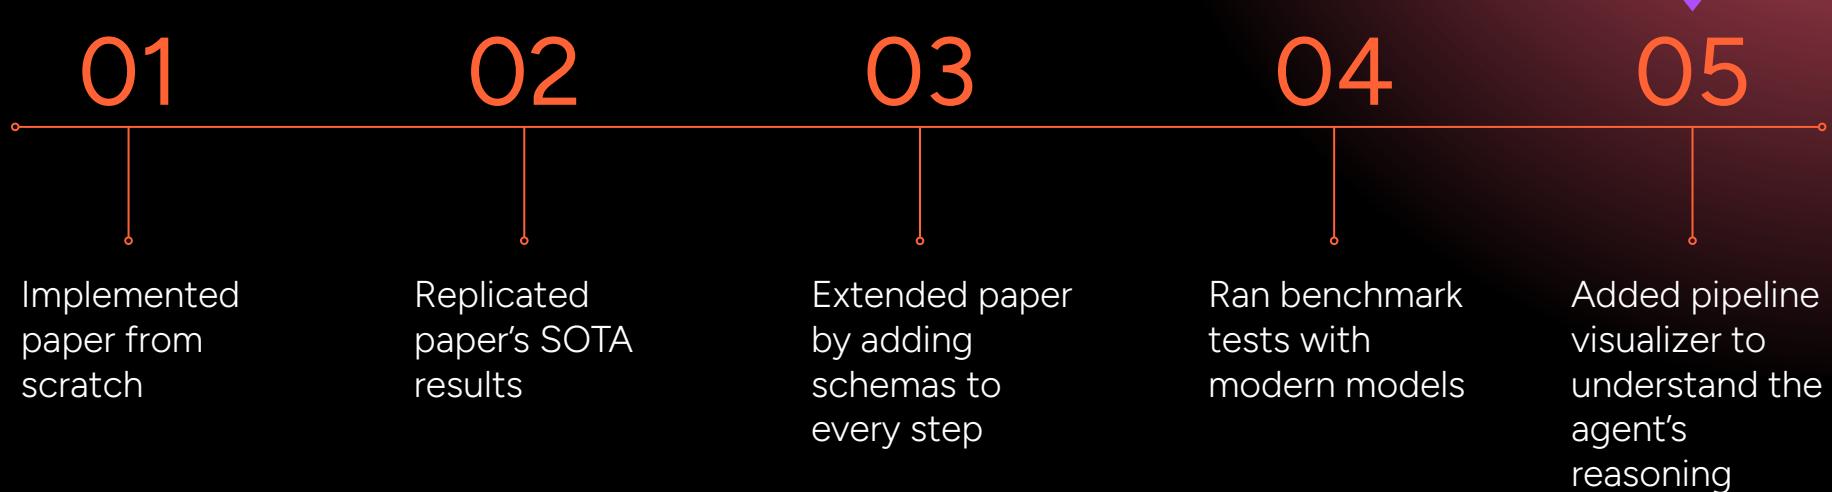
order\_by/limit missing,  
extra\_values

# TIMELINE

🔗 GitHub: [https://github.com/Jasper-256/sql\\_of\\_thought\\_recreation](https://github.com/Jasper-256/sql_of_thought_recreation)

📄 Paper: arXiv:2509.00581

We are here



# Our Implementation

## Production Ready SQL-of-Thought Framework

### Core Implementation

- Full 6-agent pipeline (LangGraph)
- Error taxonomy with 31 subtypes
- Chain-of-thought reasoning
- Guided correction loop
- Support for multiple LLM provider models

### Benchmarking

- Spider, Spider-Realistic, Spider-SYN
- Automatic dataset download
- SQLite execution & validation
- CSV output with all metrics

### Tech Stack

- Langgraph
- Langchain
- Open AI API
- Hugging Face Datasets

# Technical Architecture: Deep Dive

## 1. LangGraph State Management

1. TypedDict State Schema
2. Shared State
3. Conditional Edges

E.g. ( execute →  
  (needs\_correction &&  
  attempt < max) ?  
  correction\_plan : END  
)

## 3. Validation

1. Syntax: SQL parsing via sqlite3
2. Execution: Catch runtime exceptions
3. Semantic: Compare with gold results

## 2. Schema Introspection

1. SQLite PRAGMA queries used to extract database structure
2. Raw DB metadata into LLM-friendly format

## 4. Error Detection Strategy

- Two Correction Triggers:
- A. Exception-based (syntax/runtime)
  - B. Result mismatch (logical errors)
- Taxonomy Injection:  
→ JSON taxonomy loaded at init

## 5. Models used

1. GPT-5.1 (no thinking)
2. GPT-5-mini (high)
3. Claude Opus 4.5

## 4. Design Patterns

1. Modularity
2. Type Safety
3. Error Resilience
4. Observability

# OUR IMPROVEMENTS

## 1. Real-Time Visualization UI

-  Live NiceGUI-based dashboard
-  View all agent inputs/outputs in real-time
-  Inspect system & user prompts for each step
-  Track benchmark progress & metrics live

## 2. Advanced Model Support

-  Reasoning model integration (claude opus 4.5 and gpt-5-mini)
-  Configurable reasoning effort (low/medium/high)
-  Easy model switching via CLI
-  Measurable performance gains with modern LLMs

## 3. Enhanced Debugging

-  Pipeline visualizer for understanding agent reasoning
-  Visual representation of reasoning chain
-  Improved debugging abilities for developers

## 4. Architecture Improvements

-  Provided the full schema to the SQL agent so that it has all relevant context
-  Optimized architecture for better results

# Pipeline Visualizer

**SQL-of-Thought Pipeline Visualizer**

**Configuration**

Model	Mode	Dataset	Limit	Max Corrections
gpt-5-nano-2025-08...	SQL-of-Thought	spider	50	2

**Run Benchmark** **Stop**

**Progress**

29/50 0.0% 100.0% 72.4%

Completed Exact Match Valid SQL Exec Accuracy

57%

**Benchmark Items**

#0 (flight\_2) Valid Exec

**Pipeline Execution Details**

#1 (world\_1)

**Pipeline Execution Details**

```
HAVING COUNT(DISTINCT CountryName) > 2
ORDER BY LanguageCount DESC, CountryName ASC;
```

**Schema Linking**

- System Prompt**
- User Prompt (input)**

23:19:32

```
DB: world_1
FULL SCHEMA:
TABLE city ( ID INTEGER, Name char(35), CountryCode char(3), District char(20), Population
INTEGER )
PRIMARY KEY: ID
```

# Pipeline Visualizer

**SQL-of-Thought Pipeline Visualizer**

**Configuration**

Model	Mode	Dataset	Limit	Max Corrections
gpt-5-nano-2025-08...	SQL-of-Thought	spider	50	2

**Run Benchmark** **Stop**

**Progress**

29/50 0.0% 100.0% 72.4%

Completed Exact Match Valid SQL Exec Accuracy

57%

**Benchmark Items**

#0 (flight\_2) Valid Exec

**Pipeline Execution Details**

#1 (world\_1)

**Pipeline Execution Details**

```
HAVING COUNT(DISTINCT CountryName) > 2
ORDER BY LanguageCount DESC, CountryName ASC;
```

**Schema Linking**

- System Prompt**
- User Prompt (input)**

23:19:32

```
DB: world_1
FULL SCHEMA:
TABLE city ( ID INTEGER, Name char(35), CountryCode char(3), District char(20), Population
INTEGER )
PRIMARY KEY: ID
```

# RESULTS

## Metrics

92%

Highest Achieved

28%

Highest Achieved

100%

Highest Achieved

Execution  
Accuracy

Results Match

Exact  
Match

SQL Query String  
Equality

Valid  
SQL

Parses and Executes

# RESULTS

SQL-of-Thought paper's execution accuracy on Spider benchmark:

Model	SQL-of-Thought
GPT-4	72.3%
Claude Opus 3	91.59%

Our execution accuracy on Spider benchmark:

Model	Base LLM	SQL-of-Thought
GPT-5.1 (no thinking)	36/50 (72%)	36/50 (72%)
GPT-5-mini (high)	41/50 (82%)	40/50 (80%)
Claude Opus 4.5	43/50 (86%)	46/50 (92%)

(Measured on 50 random queries from Spider)

# THANK YOU

Any questions?